

Inverse Q-Learning as a Tool to Investigate Behavior and its Neural Correlates

Albert-Ludwigs-Universität Freiburg



UNI
FREIBURG



// BrainLinks
BrainTools

Joschka Boedecker jboedeck@cs.uni-freiburg.de

Hao Zhu hao.zhu.10015@gmail.com

September 26th, 2023

Main Questions for Today

How do we explain goal-directed animal behavior given that we often see objectively non-optimal behavior? Which factors contribute? *What are the animals optimizing for?*

Response-Preparation Task (simplified):



lever press


vibration cue

reward

- Press lever until cue (vibration) occurs (delay **1.6s**).
- After the cue, the rat has **0.6s** to release the lever
- If successful, the rat gets a treat.

Images courtesy of the Diester Lab

Reinforcement Learning in a Nutshell



$$p(s', r|s, a) = \Pr\{S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a\}$$

Goal: find policy that maximizes expected long-term reward

States



3,1	3,2	3,3	3,4
2,1	2,2	2,3	2,4
1,1	1,2	1,3	1,4

States in Autonomous Driving Application

20 features total:


Max. 6 potential vehicles surrounding the RL agent

3 features per vehicle → 18


- Relative Distance
- Relative Velocity
- Relative Angle

2 features describing the RL agent


- Velocity
- Relative Angle




Actions



Transition Probabilities



Rewards



- Stochastic vs deterministic
- Dense vs sparse (delayed)
- Magnitude

Rewards

-0.04	-0.04	-0.04	+1
-0.04		-0.04	-1
	-0.04	-0.04	-0.04

- Stochastic vs deterministic
- Dense vs sparse (delayed)
- Magnitude

Markov Decision Process

A finite Markov Decision Process (MDP) is a 4-tuple $\langle \mathcal{S}, \mathcal{A}, p, \mathcal{R} \rangle$, where

- \mathcal{S} is a finite number of states,
- \mathcal{A} is a finite number of actions,
- p is the transition probability function $p : \mathcal{S} \times \mathcal{R} \times \mathcal{S} \times \mathcal{A} \mapsto [0, 1]$
- \mathcal{R} is a finite set of scalar rewards. We can then define expected reward

$$r(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$

Markov Property:

$$\Pr\{S_{t+1}, R_{t+1} | S_t, A_t\} = \Pr\{S_{t+1}, R_{t+1} | S_t, A_t, \dots, S_0, A_0\}$$

The future is independent of the past given the present.

Policy and overall Goal

Policy determines action selection for each state:


- Stochastic: $\pi(a|s) = \Pr[A_t = a|S_t = s]$
- Deterministic: $\pi(s) = a$

Goal for an RL agent in an MDP: find a policy that maximizes the expected return, i.e. the (discounted) cumulative reward: G_t

- Finite horizon: $\arg \max_{\pi} \mathbb{E}[R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_T]$
- Infinite horizon: $\arg \max_{\pi} \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots] = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}]$

With discount $\gamma \in [0, 1]$ preventing infinite returns (converging geometric series)


Policy in MDP example



Actions:



Probability of executing action successfully: 0.8



Rewards:

-0.04 / step


Question



How will the policy change if we change the immediate reward to -2 instead of -0.04?




Policy in MDP example (changed rewards)



Actions:



Probability of executing action successfully: 0.8



Rewards:

-2 / step

Value Function and Action-Value Function

Value Function $v_\pi(s)$ is the expected return when starting in s and following π :

$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \middle| S_t = s \right]$$

Action-Value Function q_π is the expected return when starting in s , taking action a and following π thereafter:

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a] = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \middle| S_t = s, A_t = a \right]$$

Value Function Example

$v_\pi(s)$ for immediate reward of -0.04, discount of 1:

0.812	0.868	0.918	
0.762		0.660	
	0.655	0.611	0.388

Bellman Optimality Equation

- A **Bellman Equation** expresses a relationship between the value of a state and the values of its successor states
- The **Bellman Optimality Equation** expresses that the value of a state under the optimal policy π_* must equal the expected return for the best action in that state

$$\begin{aligned}v_*(s) &= \max_a q_{\pi_*}(s, a) \\&= \max_a \mathbb{E}_{\pi_*}[R_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a] \\&= \max_a \mathbb{E}_{\pi_*}[R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a] \\&= \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]\end{aligned}$$

- The value function v_* is the **unique solution** to the Bellman Optimality Equation

Bellman Equation for v_*

$$v_*(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]$$

We no longer need to search over all policies, only over all actions recursively!

Value Iteration Algorithm

An algorithm that turns the **Bellman Equation** into an **iterative update** to solve a given MDP

Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

$$\Delta \leftarrow 0$$

Loop for each $s \in \mathcal{S}$:

$$v \leftarrow V(s)$$

$$V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$$

$$\Delta \leftarrow \max(\Delta, |v - V(s)|)$$

until $\Delta < \theta$

Output a deterministic policy, $\pi \approx \pi_*$, such that

$$\pi(s) = \arg \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$$

Bellman Optimality Equation

Bellman Optimality Equation for Q-values

$$q_*(s, a) = \sum_{s', r} p(s', r | s, a) [r + \gamma \max_{a'} q_*(s', a')]$$

We no longer need to search over all policies, only over all actions recursively!

Action selection: $\pi(s) \doteq \arg \max q_\pi(s, a)$

Calculating optimal Q-values: Q-Learning

[Watkins, 1989]



UNI
FREIBURG

Q-learning

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

 until S is terminal


Figure from: [Sutton & Barto, 2018]

Standard vs Inverse RL

Standard RL:

estimate optimal **policy**
from state, action, and reward sequences

$$(s_t, a_t, s_{t+1}, r_{t+1}, \dots, s_{t+n}, r_{t+n})$$



$$\pi_\theta(s)$$

Environment



Rewards



RL



Behavior

Inverse RL:

estimate unknown **reward** function
from state and action sequences

$$(s_t, a_t, s_{t+1}, a_{t+1}, \dots, s_{t+n})$$



$$r_\psi(s, a)$$

Environment



Rewards




IRL



Behavior

Can we learn rewards from behavioral data?



Example credits:
Sergey Levine

The problem is underspecified: many reward functions would explain the behavior!

Can we learn rewards from behavioral data?

The problem is underspecified: many reward functions would explain the behavior!

Idea:

Account for uncertainty in the reward function by assuming a probabilistic behavior model that keeps action distribution in the policy as broad (non-committed) as possible



Maximum Entropy Inverse Reinforcement Learning

Problem: Needs to solve a full RL problem to convergence in the inner loop!

Deep Inverse Q-Learning (with Constraints)

[G. Kalweit, M. Huegle, M. Werling, J. Boedecker, NeurIPS, 2020]



UNI
FREIBURG

Joint work with:



Gabriel
Kalweit



Maria
Kalweit



Moritz
Werling

Deep Inverse Q-Learning

[G. Kalweit, M. Huegle, M. Werling, J. Boedecker, NeurIPS, 2020]



UNI
FREIBURG

Probabilistic behavior assumption for the expert (here for two actions, a and b):

$$\frac{\exp(Q^*(s, a))}{\exp(Q^*(s, a)) + \exp(Q^*(s, b))} = \pi^{\mathcal{E}}(a|s) \quad \text{and} \quad \frac{\exp(Q^*(s, b))}{\exp(Q^*(s, a)) + \exp(Q^*(s, b))} = \pi^{\mathcal{E}}(b|s)$$

$$\rightarrow \exp(Q^*(s, a)) + \exp(Q^*(s, b)) = \frac{\exp(Q^*(s, a))}{\pi^{\mathcal{E}}(a|s)} = \frac{\exp(Q^*(s, b))}{\pi^{\mathcal{E}}(b|s)}$$

$$\rightarrow \exp(Q^*(s, a)) = \frac{\pi^{\mathcal{E}}(a|s)}{\pi^{\mathcal{E}}(b|s)} \exp(Q^*(s, b))$$

Taking logs: $Q^*(s, a) = Q^*(s, b) + \log(\pi^{\mathcal{E}}(a|s)) - \log(\pi^{\mathcal{E}}(b|s))$

Deep Inverse Q-Learning

[G. Kalweit, M. Huegle, M. Werling, J. Boedecker, NeurIPS, 2020]



$$Q^*(s, a) = Q^*(s, b) + \log(\pi^{\mathcal{E}}(a|s)) - \log(\pi^{\mathcal{E}}(b|s))$$

Using:

$$Q^*(s, a) = r(s, a) + \gamma \max_{a'} \mathbf{E}_{s' \sim \mathcal{M}(s, a, s')} [Q^*(s', a')]$$

and replacing the Q-values above to solve for the immediate reward leads to:

$$\begin{aligned} r(s, a) &= \log(\pi^{\mathcal{E}}(a|s)) - \gamma \max_{a'} \mathbf{E}_{s' \sim \mathcal{M}(s, a, s')} [Q^*(s', a')] + r(s, b) \\ &\quad - \log(\pi^{\mathcal{E}}(b|s)) + \gamma \max_{b'} \mathbf{E}_{s' \sim \mathcal{M}(s, b, s')} [Q^*(s', b')]. \end{aligned}$$

Intuitively: immediate reward encodes the local probability of action a while also ensuring the probability of the maximizing next action a' under Q-learning

Deep Inverse Q-Learning

[G. Kalweit, M. Huegle, M. Werling, J. Boedecker, NeurIPS, 2020]



Defining: $\eta_s^a := \log(\pi^\mathcal{E}(a|s)) - \gamma \max_{a'} \mathbf{E}_{s' \sim \mathcal{M}(s, a, s')} [Q^*(s', a')]$

After some manipulation, the reward for n actions can be derived as:

$$r(s, a) = \eta_s^a + \frac{1}{n-1} \sum_{b \in \mathcal{A}_{\bar{a}}} r(s, b) - \eta_s^b.$$

This leads to three novel algorithms:

Inverse Action-Value
Iteration
IAVI

discrete state-spaces, model-based, non-linear rewards

Tabular (Constrained)
Inverse Q-Learning
(C)IQL

discrete state-spaces, sampling-based, non-linear rewards

Deep (Constrained)
Inverse Q-Learning
D(C)IQL



continuous state-spaces, sampling-based, non-linear rewards

Inverse Q-Learning: Results


[G. Kalweit, M. Huegle, M. Werling, J. Boedecker, NeurIPS, 2020]



Toy-Benchmark:
Objectworld



	IAVI	IQL	MaxEnt	Single Step MaxEnt
EVD	0.09 ± 0.00	1.47 ± 0.14	11.58 ± 0.00	4.33 ± 0.00
Runtime	$0.03 \pm 0.0 \text{ h}$	$0.35 \pm 0.0 \text{ h}$	$8.08 \pm 1.0 \text{ h}$	$12.2 \pm 0.8 \text{ h}$



Expert Demonstrations
on US Highway

Closed-form Inverse RL for Neural Decoding

[Kalweit et al., ICML Comp Bio WS, 2021]



UNI
FREIBURG

Joint work with:



// BrainLinks
BrainTools



Gabriel
Kalweit



Maria
Kalweit
(Hügle)




Ilka
Diester



Mansour
Alyahyay

Approach Summary






① Infer intrinsic reward from recorded trajectories via inverse Q-learning


② Map neural signals to intrinsic reward

③ Predict behavior via intrinsic reward function and Q-learning

Training



Reward Estimation via IAVI



IAVI returns a **scalar reward function** precisely **encoding the recorded behavior** as an intermediate result, which can then be used for neural decoding

Per-Trial Release Behavior Prediction


- Compare **NeuRL (ours)** to a random controller, logistic regression (**LR**) and non-linear classification via neural-networks (**NNC**)
- For NeuRL and NNC, **optimize hyperparameters** with random search with 500 sampled configurations each
- Consider **release prediction** in a trial if controller assigns a **probability of $> \epsilon$** (here $\epsilon = 0.6$) to action release in a given time step
- Results evaluated using **10-fold cross-validation** on a **test set**

	Rat 1			Rat 2		
	Exact Match	Near 1 Match	Near 2 Match	Exact Match	Near 1 Match	Near 2 Match
NeuRL	0.36(± 0.11)	0.49(± 0.13)	0.59(± 0.09)	0.44(± 0.09)	0.62(± 0.06)	0.70(± 0.11)
NNC	0.21(± 0.09)	0.28(± 0.12)	0.37(± 0.17)	0.34(± 0.10)	0.46(± 0.09)	0.52(± 0.10)
LR	0.15(± 0.07)	0.19(± 0.10)	0.29(± 0.08)	0.33(± 0.09)	0.41(± 0.08)	0.47(± 0.10)
Random	0.04(± 0.07)	0.2(± 0.13)	0.29(± 0.15)	0.12(± 0.06)	0.38(± 0.07)	0.46(± 0.10)


Visualization of Latent Embeddings



UNI
FREIBURG




Testing








Inhibition Simulation I



UNI
FREIBURG




Inhibition Simulation II



Outlook: IQL with Multiple Intentions




Recent (exciting!) extension by Hao Zhu




- ▼ Initial State
- Water Resource
- Food Resource
- Expert Trajectories



State-value





'Hungry'



'Thirsty'

Performance of Reward Function Estimation



		LV-IAVI	LV-IQL	IAVI	IQL
EVD	'Hungry'	0.00 ± 0.00	0.00 ± 0.00	21.42 ± 0.00	21.38 ± 0.00
	'Thirsty'	0.00 ± 0.00	0.00 ± 0.00	38.00 ± 0.00	38.05 ± 0.00

Exercises

Colab Notebook:


https://colab.research.google.com/drive/1YbHB0V1JQ5e_0T5zIR-nmRwmNOILY6v-?usp=sharing

Complete and play around with Value Iteration and Q-Learning for different tasks

Further Resources



UNI
FREIBURG



Standard RL text book (very accessible, free PDF):

<http://incompleteideas.net/book/RLbook2020.pdf>

The image shows a screenshot of a Coursera course page. At the top, there's a large green "RL" logo with a small robot icon next to it. Below the logo, the text "University of Alberta" and "Reinforcement Learning" is visible. A brief description follows: "Skills you'll gain: Machine Learning, Reinforcement Learning, Algorithms, Human Learning, Machine Learning Algorithms, Applie...". Underneath, a star rating of "4.7" with "(3.3k reviews)" is shown, along with the text "Intermediate · Specialization · 3 - 6 Months".

Nice 4-part course on coursera:

<https://www.coursera.org/specializations/reinforcement-learning>

Further Resources



UNI
FREIBURG

The screenshot shows the Google DeepMind website. At the top, there's a blue header bar with the Google DeepMind logo and navigation links: Research, Blog, Impact, Safety & Ethics, About, and Careers. Below the header is a large red section containing a portrait of a man (David Silver) on the left and text on the right. The text reads "Learning resources" and "Introduction to Reinforcement Learning with David Silver".

Very good “classic” RL course:

<https://www.deepmind.com/learning-resources/introduction-to-reinforcement-learning-with-david-silver>

CS285

CALENDAR RESOURCES SYLLABUS STAFF MENU ≡

CS 285 at UC Berkeley

Deep Reinforcement Learning

Lectures: Mon/Wed 5-6:30 p.m., Wheeler 212

NOTE: Please use the Ed link [here](#) instead of in the slides.

Lecture recordings from the current (Fall 2023) offering of the course: watch [here](#)

Looking for deep RL course materials from past years?

Recordings of lectures from Fall 2022 are [here](#), and materials from previous offerings are [here](#).

Email all staff (**preferred**): cs285-staff-fa2023@lists.eecs.berkeley.edu



Instructor **Sergey Levine**

slevine@eecs.berkeley.edu

Office Hours: After lecture



Head GSI **Kyle Stachowicz**

kstachowicz@berkeley.edu

Office Hours: Thursday 5PM–6PM
(BWW Room 1204)

Very comprehensive (Deep) RL course at UC Berkeley:

<http://rail.eecs.berkeley.edu/deeprlcourse/>